

When Structured Evidence Completion Helps—and When It Does Not: A Verified Study of HyperGranular Retrieval-Augmented Generation

Anonymous ACL submission

Abstract

Multi-hop retrieval-augmented generation must reconcile evidence completeness with a fixed Top- k context: a relevance ranking can omit a complementary item, whereas expansion can displace evidence that was already useful. We study HyperGranular-RAG, a Gold-free evidence-completion framework that combines adaptive granular balls, query-aware facet hyperedges, and bounded insertion after a protected dense prefix. Across three frozen, closed-candidate HotpotQA and MuSiQue boundaries using Qwen2.5-1.5B-Instruct, the static system repeatedly improved answer F1 over a historical MiniLM dense backbone by +0.01478 [0.00020, 0.02988], +0.01140 [0.00450, 0.01835], and +0.01357 [0.00491, 0.02233]. Removing facet hyperedges reduced performance within the frozen MiniLM system (+0.01336 [0.00341, 0.02343] for FULL minus NoFacet). The boundary changes under a stronger retriever. The original full system underperformed a pre-specified BGE baseline (−0.03998 [−0.05393, −0.02621]), while both cross-space sidecar integration (−0.00256 [−0.00998, 0.00458]) and BGE-native reconstruction (−0.003046 [−0.006880, 0.000631]) were inconclusive relative to BGE. Protected placement was supported only when the inserted set was held identical; its independent BGE-native contrast was inconclusive. Transfer to one official Gemma mobile-QAT configuration was also inconclusive. Post-decision audits further show that small net increases in Gold evidence need not improve answers. Independent verification, Gold isolation, deterministic reruns, and complete reporting of positive, negative, inconclusive, and undefined outcomes support a bounded conclusion: tested structural completion can recover evidence gaps in a compact historical backbone, but no incremental answer-quality value has been established over strong BGE.

1 Introduction

Retrieval-augmented generation (RAG) gives language models access to non-parametric evidence at inference time (Lewis et al., 2020; Guu et al., 2020). Dense passage retrieval provides an effective learned relevance ordering (Karpukhin et al., 2020), but multi-hop questions introduce a joint-coverage requirement: the selected context must contain multiple pieces of evidence, not merely high-scoring passages considered independently. HotpotQA and MuSiQue make this requirement explicit through multi-hop questions and supporting-evidence annotations (Yang et al., 2018; Trivedi et al., 2022).

A fixed Top- k budget creates a specific tension. If individually similar units occupy much of the ranking, a complementary bridge may remain outside the generator context. Expanding the ranking can recover that bridge, but every insertion can displace an item already selected by the base retriever. Retrieved evidence is also not uniformly useful to a generator: position, redundancy, prompt construction, and the generator itself affect whether a retrieved item changes the answer (Liu et al., 2024; Ram et al., 2023). Evidence coverage, ranking placement, and answer utility must therefore be measured separately.

HyperGranular-RAG addresses this constrained completion problem. It organizes sentence units into adaptive granular balls, identifies cross-ball relations through query-aware facet hyperedges, and inserts at most four eligible units after a protected Dense Top-10 prefix while preserving an effective Top-20. The design is Gold-free: candidate construction and ranking do not use answers, supporting-fact labels, or downstream answer scores. Granular balls provide adaptive local organization, while hyperedges represent relations involving more than pairwise adjacency (Xia et al., 2019; Feng et al., 2019).

This paper asks not only whether the method helps, but where its incremental value ends. We evaluate the static system over frozen HotpotQA and MuSiQue closed-candidate boundaries; isolate component contrasts; test one additional generator configuration; bind a pre-specified BGE large-en-v1.5 strong dense baseline; attach the frozen MiniLM-HyperGranular-RAG system as a BGE sidecar; and reconstruct the structural expansion natively in BGE space. Development choices are kept separate from confirmation, and all formal outcomes are independently verified.

The results support three bounded positive claims. First, static HyperGranular-RAG repeatedly improves answer F1 over the historical MiniLM Dense backbone under the frozen Qwen boundaries. Second, facet hyperedges have incremental value in the frozen MiniLM system. Third, when the Stage4I candidate set is held identical, protected placement is better than direct unprotected placement. The same evidence establishes a strong-retriever boundary: original FULL is clearly below BGE, and neither the cross-space sidecar nor BGE-native reconstruction establishes an answer-quality increment over BGE. We do not claim state-of-the-art performance, universal retrieval improvement, full-wiki validity, broad cross-generator robustness, or an independently isolated granular-ball effect.

Our contributions are:

- a protected, high-order evidence-completion construction under a fixed Top- k budget;
- a multi-boundary evaluation that retains compact-backbone gains and strong-retriever limits in one claim system;
- component evidence that distinguishes candidate construction from placement;
- direct evidence that Gold-count changes do not substitute for end-to-end answer evaluation; and
- a traceable protocol with Gold isolation, frozen artifacts, deterministic reruns, and independent verification.

2 Related Work

Dense and multi-hop retrieval. Dense passage retrieval learns dual-encoder representations for efficient passage ranking (Karpukhin et al., 2020).

Multi-hop variants extend retrieval from independent passages to evidence chains. Multi-Hop Dense Retrieval conditions later retrieval on preceding evidence (Xiong et al., 2021), while BeamDR uses beam search over dense reasoning paths (Zhao et al., 2021). PullNet iteratively expands from retrieved entities over text and knowledge bases (Sun et al., 2019), and IRCOT interleaves retrieval with chain-of-thought reasoning (Trivedi et al., 2023). HyperGranular-RAG is not an iterative question-rewriting method. Within each frozen closed candidate set, it performs a single bounded completion of an existing ranking.

Graph, hypergraph, and hierarchical retrieval.

Structure-aware retrieval represents relations that a flat relevance ordering can miss. RAPTOR retrieves from a tree of recursively summarized text (Sarathi et al., 2024). HippoRAG combines a knowledge graph with personalized propagation (Gutiérrez et al., 2024), and GNN-RAG uses graph-neural retrieval over knowledge graphs (Mavromatis and Karypis, 2025). GraphRAG targets query-focused summarization using graph communities (Edge et al., 2024). HyperGranular-RAG differs in its evidence unit, relation, and budget contract: adaptive balls organize local sentence units; facets define query-conditioned cross-ball hyperedges; and insertion preserves a dense prefix instead of replacing the complete ranking. We do not directly compare against these graph systems.

Evidence expansion and generation. Retrieval-enhanced models can expose external evidence, but the generator still consumes a finite context. RECOMP studies selective and compressed augmentation (Xu et al., 2024), while long-context studies show that evidence position affects model use even when the information is present (Liu et al., 2024). Fusion-in-Decoder demonstrates the importance of how multiple passages are consumed (Izacard and Grave, 2021), and in-context RALM shows that retrieved text can condition language models without parameter updates (Ram et al., 2023). We therefore use paired answer F1 as the primary endpoint and treat retrieval coverage as supporting evidence.

Evaluation and negative evidence. KILT emphasizes provenance-aware evaluation in knowledge-intensive tasks (Petroni et al., 2021). Reproducibility guidance motivates explicit dependencies, artifacts, and reporting boundaries

(Pineau et al., 2021). Statistical uncertainty is not resolved by relabelling an interval crossing zero as “no effect” or equivalence (Amrhein et al., 2019). We retain four evidence states: supported, negative, inconclusive, and not fairly defined.

3 HyperGranular-RAG

3.1 Problem Formulation

For a query q , let $U_q = \{u_i\}$ be the closed candidate set of sentence units. A frozen encoder maps the query and each title + “. ” + sentence unit to L2-normalized vectors. The base relevance score is cosine similarity,

$$s(q, u_i) = \hat{e}_q^\top \hat{e}_i. \quad (1)$$

Ties are resolved by ascending `unit_id`. The Dense baseline returns the first $K_q = \min(20, |U_q|)$ unique units. HyperGranular-RAG constructs an additional ordered set of eligible units but must return the same effective K_q . No answer, Gold supporting unit, or generator output enters this construction.

3.2 Adaptive Granular Balls

Each query’s candidate embeddings begin as one root set. For a current ball B , the implementation computes a normalized centroid, cosine distances to that centroid, a radius equal to the maximum distance, and compactness equal to one minus the mean distance. A ball is eligible for a deterministic two-seed split when depth is below six, the set can support two children of at least two units, and its size or radius crosses the frozen split condition. Each unit is assigned to the closer seed; the split is retained only if both children meet the minimum size. The resulting partition adapts local granularity to candidate geometry rather than imposing a fixed cluster count.

3.3 Query-Aware Facet Hyperedges

Each ball aggregates content terms from titles and sentences after frozen ASCII-alphanumeric and stop-word rules. Its facet terms are the intersection between ball terms and query terms. The two balls with highest query-to-centroid cosine become seeds. A non-seed ball can form an expansion hyperedge only if it introduces a new query term and passes frozen gates on ball size, ball score or seed proximity, units per new term, and shared-term redundancy.

The frozen facet score combines the new-facet ratio, total-facet ratio, query-to-ball score, seed diversity, redundancy, and size. Hyperedges below 0.10 are rejected; at most two edge/ball structures are selected. Units inside an accepted ball remain ordered by the original MiniLM cosine and `unit_id`; the facet score does not alter the unit score. Thus, the hyperedge selects candidate regions rather than learning a new dense ranker.

3.4 Protected Bounded Insertion

The compact variant protects the first ten unique Dense units. HyperGranular-RAG units whose frozen cosine is at least 0.1957079917192459 are considered in their frozen order. At most four candidates absent from the protected prefix are inserted immediately after that prefix. Remaining Dense units are appended in their original order, followed only if necessary by unfilled expansion units, and the result is truncated to K_q .

Protection and efficacy are separate claims. Protection limits displacement, but it cannot make a non-complementary candidate useful. This distinction motivates the placement and BGE comparisons below.

3.5 Compact, Sidecar, and Native Variants

The *compact variant* uses MiniLM embeddings for the base ranking, ball geometry, facet gates, and unit ordering. The *cross-space sidecar* keeps BGE Top-20 as the primary ranking while using the frozen MiniLM-HyperGranular-RAG structure to propose eligible candidates; no score fusion is introduced. The *BGE-native variant* reconstructs candidate geometry and structural selection in BGE space under one development-selected, then frozen, configuration. These variants test different scientific questions and are not pooled.

4 Experimental Protocol

4.1 Datasets and Frozen Boundaries

HotpotQA provides diverse multi-hop questions and supporting facts (Yang et al., 2018); MuSiQue composes single-hop questions to reduce shortcut solutions (Trivedi et al., 2022). All experiments use benchmark-supplied closed candidate sets. They do not query a full Wikipedia index.

The compact program includes a 1,000-query HotpotQA boundary, a 3,000-query MuSiQue boundary, and a separate 1,000 HotpotQA plus

HyperGranular-RAG: protected high-order evidence completion under a fixed Top-k budget




Figure 1: **HyperGranular-RAG method overview.** The compact system organizes locally related evidence as adaptive granular balls, links query-relevant facets across balls with high-order hyperedges, and inserts bounded candidates after a protected prefix under the same Top- k budget. Dashed cross-space and BGE-native variants are evaluated configurations, not validated improvements over BGE.

1,500 MuSiQue component boundary. The cross-space sidecar and BGE-native confirmation studies each use their own frozen, zero-overlap 1,000 HotpotQA plus 1,500 MuSiQue boundary. Selection is deterministic and identity audited.

4.2 Retrievers and Generators

The historical compact backbone is the frozen MiniLM dense retriever. BGE large-en-v1.5 is the pre-specified strong dense baseline; the broader BGE family is documented by Xiao et al. (2023), while the exact model revision is bound in repository manifests.

The main generator is Qwen2.5-1.5B-Instruct FP16 with a fixed short-answer RAG prompt and a 4,096-token input cap. One additional transfer evaluation uses google/gemma-4-E2B-it-qat-mobile-transformers in its official mobile-QAT format. The Qwen-Gemma comparison is deployment-bound: architecture, quantization/numerical format, and runtime effects are not isolated. It cannot support a pure model-architecture ranking or a claim about all devices.

4.3 Development-Confirmation Separation

All method and decision parameters are frozen before the corresponding confirmation boundary is opened. The BGE-native development study selected configuration C10 with a development answer-F1 difference of +0.003143. That value is selection evidence only. It is neither pooled with confirmation nor presented as a confirmed positive

result. Automatic BGE-native parameter search stopped after the pre-registered confirmation outcome.

4.4 Metrics and Statistical Decisions

The primary endpoint is paired query-level answer-F1 difference. Answer exact match (EM), retrieval complete recall at 20 (CR@20), and evidence recall at 20 (ER@20) are supporting endpoints. Frozen comparisons use 10,000 paired bootstrap samples and stage-specific pre-specified decision rules. Dataset-joint comparisons assign equal weight to HotpotQA and MuSiQue. An interval crossing zero is INCONCLUSIVE; no equivalence margin or formal equivalence test was specified. Sample sizes define available resource and interval precision, not a formal power guarantee.

4.5 Gold Isolation and Independent Verification

Retrieval, prompt construction, and generation finish before Gold access. Formal evaluators verify Gold identity before parsing it. Independent verifiers reconstruct query order, metrics, bootstrap summaries, decisions, and artifact identities. Main/rerun transactions are either byte-identical in full or use a pre-hash deterministic subset whose selected outputs match the corresponding main outputs exactly. All final scientific decisions and manifests were frozen before manuscript construction.

Boundary	Δ F1 [95% CI]	State
HotpotQA	+0.01478 [0.00020, 0.02988]	Supported
MuSiQue	+0.01140 [0.00450, 0.01835]	Supported
Joint	+0.01357 [0.00491, 0.02233]	Supported

Table 1: Repeated FULL minus DENSE results under the historical MiniLM backbone, Qwen generator, closed candidates, and Top-20 budget.

Contrast	Δ F1 [95% CI]	State
MiniLM facet	+0.01336 [0.00341, 0.02343]	Supported
Sidecar placement	+0.01122 [0.00129, 0.02104]	Supported
Native placement	+0.00337 [−0.00129, 0.00803]	Inconcl.
Native facet	+0.00063 [−0.00422, 0.00536]	Inconcl.

Table 2: Component evidence. “Native” denotes the BGE-native confirmation boundary; effects are not transported between semantic spaces.

5 Results

5.1 Repeated Gains over Compact Dense Retrieval

Table 1 and Figure 2 show three frozen Qwen comparisons. On 1,000 HotpotQA queries, FULL minus DENSE answer F1 is +0.01478 [0.00020, 0.02988]. On 3,000 MuSiQue queries, the difference is +0.01140 [0.00450, 0.01835]. The separate joint component boundary again supports FULL minus DENSE with a dataset-equal-weight difference of +0.01357 [0.00491, 0.02233].

The effects are small but directionally repeated. Their scope is the historical MiniLM backbone, Qwen generator, fixed prompt, closed candidates, and Top-20 budget. They do not establish advantage over modern strong retrievers.

5.2 Component Evidence

Within the frozen MiniLM system, FULL minus NoFacet is +0.01336 [0.00341, 0.02343], supporting an incremental role for facet hyperedges. FULL minus NoProtection is +0.00354 [−0.00675, 0.01389], so the independent contribution of protection is inconclusive on that boundary. A flat-unit granular-ball contrast is NOT_FAIRLY_DEFINED: ball seeds, split gates, edge budgets, and radii have no unique matched flat equivalents. The absence of this contrast is not evidence for or against a granular-ball effect.

In the cross-space sidecar, the inserted set is identical between protected and unprotected conditions. PROTECTED minus Unprotected is +0.01122 [0.00129, 0.02104], which supports the placement rule within that specific sidecar. However, PROTECTED minus NoFacet is −0.00743 [−0.01616, 0.00132], inconclusive. In BGE-native space, PROTECTED minus Unprotected is +0.003367 [−0.001285, 0.008026], and PROTECTED minus NoFacet is +0.000631 [−0.004220, 0.005365]; both are inconclusive. Component effects are therefore system-dependent.

5.3 Strong-Dense Baseline Boundary

The pre-specified BGE baseline outperforms original FULL. The dataset-equal-weight FULL minus BGE difference is −0.03998 [−0.05393, −0.02621]. This negative result is qualitatively different from an uncertain contrast: the complete interval lies below zero. It rules out the claim that the original compact-backbone HyperGranular-RAG system is competitive with this strong BGE retriever on the frozen boundary.

5.4 Cross-Space Sidecar Evaluation

The cross-space sidecar asks whether the frozen MiniLM structural expansion can add complementary candidates while BGE retains the primary ranking. Its core PROTECTED minus BGE difference is −0.00256 [−0.00998, 0.00458]. Because the interval crosses zero, the result does not establish gain, harm, or equivalence. The positive protected-placement contrast does not change this efficacy result: better placement of a fixed inserted set is not proof that the set improves BGE.

5.5 BGE-Native Reconstruction

The BGE-native study removes the cross-space mismatch by rebuilding candidate geometry and facets in BGE space. Confirmation remains inconclusive. The dataset-equal-weight PROTECTED minus BGE answer-F1 difference is −0.003046 [−0.006880, 0.000631]; the EM difference is −0.003667 [−0.007667, 0.000004]. Dataset-specific F1 differences are −0.002907 [−0.008682, 0.002652] for HotpotQA and −0.003184 [−0.008486, 0.001863] for MuSiQue. Both point estimates are slightly negative, but all intervals cross zero. Thus the tested BGE-native reconstruction establishes neither improvement, damage, nor equivalence.

The confirmation result must not be merged with the positive C10 development difference. The development value selected the configuration; the confirmation boundary evaluated it.




Figure 2: **Frozen answer-F1 effect sizes.** Points are paired answer-F1 differences and bars are pre-specified 95% bootstrap confidence intervals. The zero line is explicit. INCONCLUSIVE means that the interval crosses zero; it is not an equivalence claim.

Boundary	Δ F1 [95% CI]	State
Gemma transfer	+0.00516 [-0.00262, 0.01295]	Inconcl.
Full vs. BGE	-0.03998 [-0.05393, -0.02621]	Negative
Sidecar vs. BGE	-0.00256 [-0.00998, 0.00458]	Inconcl.
Native vs. BGE	-0.00305 [-0.00688, 0.00063]	Inconcl.

Table 3: Strong-retriever boundary and generator transfer. Each row uses its own frozen evaluation boundary; rows are not pooled.

5.6 Generator Transfer

Under the single pre-specified Gemma mobile-QAT configuration, the dataset-equal-weight FULL minus DENSE difference is +0.00516 [-0.00262, 0.01295], inconclusive. This is not evidence that Gemma is unsuitable for RAG, that Qwen has universally superior base capability, or that one architecture is better. The tested model formats differ, and the result is restricted to the recorded RTX 4060 Laptop 8GB, short-answer prompt, 4,096-token cap, and deployable formats.

5.7 Efficiency and Evidence Displacement

Resource values are hardware-bound descriptors rather than cross-device efficiency claims. For the Stage5A transaction, retrieval reconstruction took

8.544 seconds; BGE and BGE-native PROTECTED generation took 887.650 and 884.198 seconds, respectively; recorded GPU peak memory was 3.592 GiB.

Post-decision Gold audits clarify why coverage and answer quality diverge (Figure 3). In the BGE-native PROTECTED condition, HotpotQA adds 3 Gold units, displaces 2, and has net +1; MuSiQue adds 16, displaces 15, and also has net +1. Nevertheless, answer-F1 point estimates relative to BGE are negative. In the cross-space sidecar, HotpotQA has 21 added, 29 displaced, and net -8, whereas MuSiQue has 59 added, 50 displaced, and net +9. These counts are post-decision descriptive and cannot be used to retune the method.

6 Discussion

Why compact retrieval can benefit. The repeated MiniLM gains are consistent with a bounded interpretation: when a compact representation leaves recoverable evidence gaps, local candidate organization plus query-conditioned cross-ball expansion can expose complementary units that a flat ranking omits. The Stage4H facet contrast supports facet-mediated candidate selection




Figure 3: **Strong-retriever boundary and evidence displacement.** Panel A compares each structural method only with BGE inside its own frozen Stage4H, Stage4I, or Stage5A boundary; the groups are not pooled. Panel B reports post-decision descriptive Gold evidence added and displaced. These counts did not control selection or advancement and do not replace end-to-end answer evaluation.

within that implementation. This is a synthesis of the frozen evidence, not a universal theorem about compact encoders.

Why the increment contracts under strong BGE. BGE begins with substantially higher answer quality than the original MiniLM system. The original FULL ranking is clearly worse than BGE, and both tested extensions around BGE yield narrow intervals around small negative point estimates. A stronger representation may already rank much of the recoverable evidence highly, leaving less marginal opportunity for structural expansion. Additional units then compete with already strong evidence. The experiment supports this as a study-specific boundary, not as a claim that structure-aware retrieval can never complement strong retrievers.

Evidence addition is not evidence utility. Stage5A is especially informative: net Gold increases slightly, but answer quality does not. A Gold-labelled unit can be redundant with retained context, weakly positioned, difficult for the generator to use, or offset by a different displacement. Conversely, a non-Gold unit can affect the answer. CR/ER, Gold transitions, and answer F1 therefore answer different questions.

Placement matters but cannot substitute for efficacy. The Stage4I placement result isolates ranking position because protected and unprotected conditions contain the same inserted set. It supports preserving high-ranked evidence under a

fixed budget. However, Stage5A does not independently confirm the placement increment, and neither placement experiment establishes that PROTECTED exceeds BGE.

Facet effects depend on semantic space. Facet hyperedges help in the frozen MiniLM system, but their increments are inconclusive in cross-space and BGE-native studies. Query-to-ball geometry, seed identities, facet eligibility, and competition with the base ranking all depend on representation space. The MiniLM effect is therefore not transported to BGE.

Development is not confirmation. The positive BGE-native development result justified opening one frozen confirmation transaction. It did not license retroactive configuration search after confirmation. Continuing to tune on that boundary would convert confirmation into development. Stopping the search is a research-integrity decision, not a claim that no future structure-aware method can work.

Inconclusive is not equivalent. The sidecar, BGE-native, and generator-transfer intervals cross zero. Without a pre-specified equivalence margin and an equivalence test, these results cannot be called equal. They also do not establish harm.

7 Limitations

Evaluation is closed-candidate rather than full-wiki or open-domain; it does not measure indexing, approximate-nearest-neighbor recall, or

corpus-scale noise. The repeated positive evidence uses one Qwen generator, and the only additional generator is a single Gemma mobile-QAT configuration. Only one strong BGE backbone is tested. The BGE-native configuration family is finite and development-selected; untested variants remain unknown.

A scientifically matched flat granular-ball control was not defined, so the granular-ball structure lacks an independent ablation claim. Protected insertion has inconsistent independent evidence: Stage4I placement is supported, while Stage4H and Stage5A protection-related contrasts are inconclusive. Facet evidence is representation- and implementation-dependent.

The HotpotQA training split may have appeared in generator pretraining. The project can guarantee only that its own retrieval and evaluation respected the frozen boundaries. Post-Gold mechanism audits are descriptive, not confirmatory causal analyses. Sample sizes reflect available resources and interval precision, not a formal power design. No inconclusive result is evidence of equivalence.

8 Reproducibility and Integrity

The project binds dataset files, model revisions, configurations, scripts, rankings, predictions, prompt audits, query audits, bootstrap summaries, decisions, and final verifications by byte length and SHA-256. Stage4E through Stage5A are frozen and were not recomputed for this manuscript. Independent verifiers reconstruct numerical outputs and decisions from formal artifacts rather than trusting report text.

Every plotted value has a CSV source; SVG, PDF, 600-dpi TIFF, and PNG exports are bound in a derived-file manifest. Two consecutive builds were byte-identical. A read-only manuscript verifier checks frozen inputs, figure and table sources, citation keys, local links, required claim strings, status wording, and unresolved citation placeholders. Dataset and model licenses are recorded separately from scientific claims. Large local caches and third-party model weights are not part of the manuscript package.

9 Conclusion

HyperGranular-RAG is a protected high-order evidence-completion framework for constrained multi-hop retrieval. Across three frozen bound-

aries, it repeatedly improves a historical compact MiniLM backbone, and facet hyperedges contribute within the frozen MiniLM system. The evidence also draws a clear limit: the original method underperforms strong BGE, while neither cross-space sidecar integration nor BGE-native reconstruction establishes an incremental answer-quality gain over BGE. Protected placement can matter for a fixed candidate set, yet placement and Gold-evidence increments cannot substitute for end-to-end efficacy. The contribution is not a universal stronger retriever; it is a verified account of when structured evidence completion helps, when its marginal value contracts, and how those outcomes should be reported without hiding negative or inconclusive evidence.

Ethical Considerations

This study uses existing benchmark datasets and publicly documented model artifacts. It does not recruit human participants or collect new personal data. The analysis reports negative, inconclusive, and undefined results alongside positive findings to reduce selective-reporting risk. A venue-specific determination of whether a formal institutional ethics statement is required remains subject to human confirmation.

Data and Code Availability

The identity-exposing development repository records code, configurations, frozen derived evaluation artifacts, figure source data, and verification manifests. It is not linked in this anonymous draft. Original benchmark datasets, model weights, and large local caches are not redistributed and must be obtained from their official sources under their respective terms. An anonymous artifact snapshot and final reuse license remain pending human decisions.

Generative AI Disclosure (Draft)

Generative AI tools were used for code assistance, experiment-workflow assistance, documentation assistance, and language-editing assistance. They are not authors and do not bear academic responsibility. Human authors retain responsibility for the scientific design, evidence selection, verification, claims, and submitted text. Final wording will be reviewed against the selected venue’s policy before submission.

References

- Valentin Amrhein, Sander Greenland, and Blake McShane. 2019. [Scientists rise up against statistical significance](#). *Nature*, 567:305–307.
- Darren Edge, Ha Trinh, Newman Cheng, Joshua Bradley, Alex Chao, Apurva Mody, Steven Truitt, and Jonathan Larson. 2024. [From local to global: A graph RAG approach to query-focused summarization](#). *arXiv preprint arXiv:2404.16130*.
- Yifan Feng, Haoxuan You, Zizhao Zhang, Rongrong Ji, and Yue Gao. 2019. [Hypergraph neural networks](#). In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 3558–3565.
- Bernal Jiménez Gutiérrez, Yiheng Shu, Yu Gu, Michihiro Yasunaga, and Yu Su. 2024. [HippoRAG: Neurobiologically inspired long-term memory for large language models](#). In *Advances in Neural Information Processing Systems*, volume 37.
- Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Ming-Wei Chang. 2020. [REALM: Retrieval-augmented language model pre-training](#). In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 3929–3938.
- Gautier Izacard and Edouard Grave. 2021. [Leveraging passage retrieval with generative models for open domain question answering](#). In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics*, pages 874–880.
- Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. 2020. [Dense passage retrieval for open-domain question answering](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, pages 6769–6781.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. [Retrieval-augmented generation for knowledge-intensive NLP tasks](#). In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474.
- Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. 2024. [Lost in the middle: How language models use long contexts](#). *Transactions of the Association for Computational Linguistics*, 12:157–173.
- Costas Mavromatis and George Karypis. 2025. [GNN-RAG: Graph neural retrieval for efficient large language model reasoning on knowledge graphs](#). In *Findings of the Association for Computational Linguistics: ACL 2025*.
- Fabio Petroni, Aleksandra Piktus, Angela Fan, Patrick Lewis, Majid Yazdani, Nicola De Cao, James Thorne, Yacine Jernite, Vassilis Plachouras, Tim Rocktäschel, and Sebastian Riedel. 2021. [KILT: A benchmark for knowledge intensive language tasks](#). In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics*, pages 2523–2544.
- Joelle Pineau, Philippe Vincent-Lamarre, Koustuv Sinha, Vincent Larivière, Alina Beygelzimer, Florence d’Alché Buc, Emily Fox, and Hugo Larochelle. 2021. [Improving reproducibility in machine learning research](#). *Journal of Machine Learning Research*, 22(164):1–20.
- Ori Ram, Yoav Levine, Itay Dalmedigos, Dor Muhlgay, Amnon Shashua, Kevin Leyton-Brown, and Yoav Shoham. 2023. [In-context retrieval-augmented language models](#). *Transactions of the Association for Computational Linguistics*, 11:1316–1331.
- Parth Sarthi, Salman Abdullah, Aditi Tuli, Shubh Khanna, Anna Goldie, and Christopher D. Manning. 2024. [RAPTOR: Recursive abstractive processing for tree-organized retrieval](#). In *International Conference on Learning Representations*.
- Haitian Sun, Tania Bedrax-Weiss, and William W. Cohen. 2019. [PullNet: Open domain question answering with iterative retrieval on knowledge bases and text](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*, pages 2380–2390.
- Harsh Trivedi, Niranjan Balasubramanian, Tushar Khot, and Ashish Sabharwal. 2023. [Interleaving retrieval with chain-of-thought reasoning for knowledge-intensive multi-step questions](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*, pages 10014–10037.
- Harsh Trivedi, Tushar Khot, Ashish Hartmann, Ruskin Manku, and Ashish Sabharwal. 2022. [MuSiQue: Multihop questions via single-hop question composition](#). *Transactions of the Association for Computational Linguistics*, 10:539–554.
- Shuyin Xia, Yunsheng Liu, Xin Ding, Guoyin Wang, Hong Yu, and Yuoguo Luo. 2019. [Granular ball computing classifiers for efficient, scalable and robust learning](#). *Information Sciences*, 483:136–152.
- Shitao Xiao, Zheng Liu, Peitian Zhang, Niklas Muenighoff, Defu Lian, and Jian-Yun Nie. 2023. [C-Pack: Packaged resources to advance general chinese embedding](#). *arXiv preprint arXiv:2309.07597*.
- Wenhan Xiong, Xiang Lorraine Li, Srinivasan Iyer, Jingfei Du, Patrick Lewis, William Yang Wang, Yashar Mehdad, Wen-tau Yih, Sebastian Riedel, Douwe Kiela, and Barlas Oguz. 2021. [Answering complex open-domain questions with multi-hop dense retrieval](#). In *International Conference on Learning Representations*.

Fangyuan Xu, Weijia Shi, and Eunsol Choi. 2024. [RE-COMP: Improving retrieval-augmented LMs with compression and selective augmentation](#). In *International Conference on Learning Representations*.

Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W. Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. 2018. [HotpotQA: A dataset for diverse, explainable multi-hop question answering](#). In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2369–2380.

Chen Zhao, Chenyan Xiong, Jordan Boyd-Graber, and Hal Daumé III. 2021. [Multi-step reasoning over unstructured text with beam dense retrieval](#). In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics*, pages 4635–4641.

A Complete Evidence-State Map

B Applicability Boundary

C Efficiency and Verification Summary

D Claim–Evidence Boundary Summary




Figure 4: **Evidence-state map.** Supported, negative, inconclusive, and not-fairly-defined outcomes are kept distinct. Component findings are system-dependent and are not transported across semantic spaces without confirmation.

Item	Value	Evidential role
Stage5A retrieval reconstruction	8.544 s	Hardware-bound descriptive
Stage5A BGE generation	887.650 s	Hardware-bound descriptive
Stage5A native Protected generation	884.198 s	Hardware-bound descriptive
Stage5A GPU peak memory	3.592 GiB	Hardware-bound descriptive
Bootstrap	10,000 iterations per frozen decision	Confirmatory
Gold policy	Isolated until rankings/predictions were frozen	Integrity
Verification	Independent final verifiers passed	Integrity

Table 4: Efficiency and verification summary. Time and memory values do not support cross-device efficiency claims.

Applicability boundary of the frozen evidence




Figure 5: **Applicability boundary.** Repeated small gains are established only for the historical compact MiniLM backbone under closed-candidate Top-20 evaluation. The original system is negative against strong BGE, the tested sidecar and BGE-native extensions are inconclusive, and full-wiki/open-domain operation was not evaluated.

Claim	Evidence state	Boundary
Compact FULL exceeds DENSE on three frozen boundaries	Supported	MiniLM, Qwen, closed-candidate Top-20
Facet increment in the frozen MiniLM system	Supported	Stage4H only
Protected placement with identical sidecar insertions	Supported	Stage4I only
Original FULL exceeds strong BGE	Negative	Stage4H
Cross-space sidecar exceeds BGE	Inconclusive	Stage4I
BGE-native PROTECTED exceeds BGE	Inconclusive	Stage5A
BGE-native placement increment	Inconclusive	Stage5A
BGE-native facet increment	Inconclusive	Stage5A
Additional-generator transfer	Inconclusive	One Gemma mobile-QAT setting
Flat granular-ball control	Not fairly defined	No matched contrast

Table 5: The manuscript’s claim system. Inconclusive rows are not equivalence claims, and system-specific component evidence is not transported across representation spaces.