

A Fitting and continuation

FIT queries
Known
support
annotations

Static sets
FIT search sets
(final / frontier)

H1 / H2 / H4 / DeepSets
Static: constructed sets
Aligned: static + search sets
H4 replay: equal updates
Separately trained parameters

BGE frozen; support supervision, no reader-answer reward.

RQ1 readout: supplied-TUNE-set discrimination.

B Selection and answering

Question
+ original
corpus

Frozen BGE
Shared
Top-128
paragraphs

Dense / MMR
budget selectors

**Learned score
+ finite beam**
≤6 whole paragraphs
Feasible Dense-K6
also eligible

H4: Flat / GB / KM
Same score;
alternatives
RQ3: choices + time

**Common
renderer**
Whole
paragraphs
BGE order

Fixed FP16 3B
Qwen2.5 / P2
Greedy decoding

Answer

Every method: actual serialized
input ≤1,024 tokens.
Dense/MMR may exceed
6 paragraphs.

C Evaluation readouts (no selection feedback)

Support
annotations

RQ1: support
Coverage /
completeness

RQ2: paired support + F1
Same 128 questions

Answer
evaluation
Canonical
EM / F1

Gold answers

Exposed development: initial TUNE includes QA;
continuation DEV-SELECT excludes QA.